


One model call, typically about two seconds



v2 · n=300 · mean 2.23 s · exclusive-quantile p95. This is translation latency, not robot mission duration.